

How to distrust AI:

Jonathan' speaking notes

- (current) AI is not a person. It's a profoundly complex prediction model. It was not born. It does not have a worldview, lived experience, or intuition. NO SOUL
 - You may've heard AI is a fancy "autocorrect" or "autocomplete". This is more true than most people realize. Started as GPT 1, 2, 3... Large Language Model (statistical model) — Completing a story.
 - "How do I bake a cake?"
 - The AI might predict the next logical sentence and output: "And how do I make frosting?" or "Please help, my son's birthday is tomorrow."
 - —> The shift to conversational AI
 - **System Prompts**
 - "You are a helpful, harmless, and honest AI assistant. You provide direct, accurate answers. Do not make things up. The user will ask you a question, and you will respond."
 - Gives the AI a character to play
 - **Chat Formatting** — Like the script for a theatrical play or movie!!
 - **[SYSTEM_START]** You are a helpful assistant.
[SYSTEM_END]
 - **[USER_START]** What is the capital of France?
[USER_END]
 - **[AI_START]**

○ **Tokens**

- We see: Letters, Syllables, Words and Sentences
- AI sees tokens. Some smaller words are one token. Many are broken up. On average, one token is roughly three quarters of a word.

○ **Input:**

- **Training Data for the LLM (longterm memory)**
- Hidden prompts and context
- Your prompt

○ **Context window (short term memory)**

- Everything the trained AI sees
- The context window is the AI's working memory.
 - Every time you prompt, it begins processing from the very beginning of the conversation all over again.
 - When the context window gets more filled up, even if it's not maxed out, you can start getting worse results. Especially if there are irrelevant things in there
 - Context pollution! Separate concerns!
 - The end-of-response hook, ignore it or start a new chat for it, unless
 - 1) it's truly what you want next, AND
 - 2) the current context is important for this thinking
 - The AI loses "attention" somewhat in the middle, "forgetting" details. Put the most important parts at the beginning or end.

- Avoid unnecessary information.
- If the conversation gets so long that the context window is truly full, the AI either completely stops working (the old early behavior), or it starts ignoring early parts of the conversation. It will literally have no awareness of those early parts other than them being vaguely referenced later.
- Better results
 - Ask it to research and cite sources. Then go check those sources yourself.
 - Put the most important rules, questions, or instructions at the very top, or very bottom
 - Provide clear boundaries. use formatting like quote marks or brackets to separate your instructions from any documents or data
 - Don't assume the AI knows anything about you that you haven't told it, or stored in the setting or memory file somewhere
 - If you're looking for a certain type of analysis, it can help to identify which frameworks you want it to use.
 - For my emotional communication analysis, I asked for it to analyze based on frameworks from:
 - Nonviolent Communication (NVC)
 - The Gottman Method
 - Interpersonal Neurobiology (IPNB) and Sarah Peyton
 - Customize the AI personality and style of responses in settings. Allows you to add your own additions to

the hidden system prompt.

- Productivity

- Talk directly to the AI! Most people can speak much faster, more fluidly, and more creatively than they can type.

- I use: Wispr Flow

- Privacy

- Setting -> TURN OFF ALLOWING TRAINING
 - Or use your own local AI system, not a remote service